


TREES



So log-arithmetic. (Get it!?)

TREE




Set of nodes and edges
(references to child
nodes) without any cycles



TREES




NOT TREES



BINARY TREE

Tree in which each node has at most two children



Size 9 Height 3

root

unique topmost node of tree
that can reach all other nodes

parent ➡ **child**

node above ➡ node below

descendant

node reachable from parent
to child, grandchild, etc.

ancestor

node reachable from child to
parent, grandparent, etc.

leaf / external node

node with no children

branch / internal node

node with at least one child

size (tree)

number of nodes in the tree

height (tree)

number of edges on *longest*
downward path from *root* to leaf

height (node)

number of edges on *longest*
downward path from *node* to
leaf

depth (node)


number of edges between the
node and the root

level (node)

1 + number of edges between
the *node* and the root


COMPLETE TREE

Every level except
possibly last is
completely filled and
nodes are as far left
as possible




BALANCED TREE

All leaves are at minimum
possible depth



Balanced



Unbalanced

BINARY SEARCH TREE


Always sorted

For each node

Left children are smaller

Right children are larger

No duplicate keys (usually)



WHY USE A BST?

Fast search, insertion, deletion -
especially when balanced

Sort as you go instead of all at once

Fairly simple implementation for good
performance

WHY USE A BST?

Only allocates memory as it's needed

Doesn't have to reallocate memory to grow (like a hash table)

A NOTE ON LOG N

When discussing time and space complexity of algorithms, $\log n$ usually means $\log_2 n$

BINARY LOGARITHM

$$\log_2 n = x \leftrightarrow 2^x = n$$

the power by which 2 must be raised by to obtain n

$$\log_2 16 = 4 \quad \log_2 32 = 5 \quad \log_2 143 = 7.15987$$

WHY IS LOG N GOOD?

Imagine a binary search tree with 2^{32} nodes

2^{32} is 4,294,967,296 (that's over 4 billion!)

When searching, we only have to visit a maximum of 32 nodes to find the node containing the data we're looking for (assuming the tree is perfectly balanced)

SEARCH

```
# Call initially with node == root node
def find_recursive(key, node):
    if node is None or node.key == key:
        return node
    elif key < node.key:
        return find_recursive(key, node.left)
    else:
        return find_recursive(key, node.right)
```


INSERTION

Same as search except once you find a node without a child on the next side you're traversing, add it there.

DELETION


Three cases

No children


One child

Two children


NO CHILDREN



ONE CHILD



TWO CHILDREN



COMPLEXITY

	Average Case	Worst Case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(n)$
Insert	$O(\log n)$	$O(n)$
Delete	$O(\log n)$	$O(n)$